

Artificial Intelligence

CSE 440/EEE 333/ETE333

Chapter 5

Fall 2017

Mirza Mohammad Lutfе Elahi

Department of Electrical and Computer Engineering
North South University

Generalizing Search Problems

- So far: our search problems have assumed agent has complete control of environment
 - state does not change unless the agent (robot) changes it.
 - All we need to compute is a single path to a goal state.
- Assumption not always reasonable
 - stochastic environment (e.g., the weather, traffic accidents).
 - other agents whose interests conflict with yours
 - Problem: you might not traverse the path you are expecting.

Generalizing Search Problems

- In these cases, we need to generalize our view of search to handle state changes that are not in the control of the agent.
- One generalization yields game tree search
 - agent and some other agents.
 - The other agents are acting to maximize their profits
 - this might not have a positive effect on your profits.

More General Games

- What makes something a game?
 - there are two (or more) agents influencing state change
 - each agent has their own interests
 - e.g., goal states are different; or we assign different values to different paths/states
 - Each agent tries to alter the state so as to best benefit itself.
- What makes games hard?
 - how you should play depends on how you think the other person will play; but how they play depends on how they think you will play; so how you should play depends on how you think they think you will play; but how they play should depend on how they think you think they think you will play; ...


Games: Language/Functions

- s_0 // initial state
- **PLAYERS**(s) // who's the player in state s
- **ACTIONS**(s) // possible moves from state s
- **RESULT**(s, a) // the state after action a is taken on state s
- **TERMINAL-TEST**(s) // returns true if s is a terminal state
- **UTILITY**(s, p) // the objective function in state s for player p


Zero-Sum Games

- **Zero-Sum Games**
 - Agents have opposite utilities (values on outcomes)
 - Lets us think of a single value that one maximizes and the other minimizes
 - Adversarial, pure competition
- **General Games**
 - Agents have independent utilities (values on outcomes)
 - Cooperation, indifference, competition, and more are all possible
 - More later on non-zero-sum games

Adversarial Search



Single-Agent Trees




Value of a State

Value of a state:
The best
achievable
outcome (utility)
from that state

Non-Terminal States:


$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$



Terminal States:

$$V(s) = \text{known}$$

Adversarial Game Trees




Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:


$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$

Tic-Tac-toe Game Tree



Game Playing Problem


- Instance of the general search problem.
- States where the game has ended are called **terminal states**.
- A **utility (payoff) function** determines the value of **terminal states**, e.g. win = +1, draw = 0, lose = -1.
- In two-player games, assume one is called **MAX** (tries to maximize utility) and one is called **MIN** (tries to minimize utility).
- In the search tree, first layer is move by **MAX**, next layer by **MIN**, and alternate to terminal states.
- Each layer in the search is called a **ply**

Optimal Decision: Optimal Players

- A game not only finds best way to goal
- The other player has a say
- **Two players:** MAX and MIN
- **ACTIONS(s)** and **RESULTS(s, a)** define a game tree
- **Tic Tac Toe:** Fewer than $9!(3,62,880)$ terminal nodes
- **Chess:** over 10^{40}
- **Search tree** as theory

Adversarial Search (Minimax)

- **Deterministic, zero-sum games:**
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- **Minimax search:**
 - A state-space search tree
 - Players alternate turns
 - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



Minimax

- General method for determining optimal move.
- Generate complete game tree down to terminal states.
- Compute utility of each node bottom up from leaves toward root.
- At each MAX node, pick the move with maximum utility.
- At each MIN node, pick the move with minimum utility (assumes opponent always acts correctly to minimize utility).
- When reach the root, optimal move is determined.

Minimax algorithm

```
function MINIMAX-DECISION(state) returns an action  
  inputs: state, current state in game  
  return the a in ACTIONS(state) maximizing MIN-VALUE(RESULT(a, state))
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for a, s in SUCCESSORS(state) do  $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$   
  return v
```


```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for a, s in SUCCESSORS(state) do  $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$   
  return v
```

Minimax

- Small 2-ply game

MAX


MIN



$$\triangle = \text{MAX}; \nabla = \text{MIN}$$

Minimax

- Picking my best move against your best move






$minimax(s) =$

$$\begin{cases} utility(s) & \text{if } terminal(s) \\ \max_{a \in action(s)} minimax(result(s, a)) & \text{if } player(s) = MAX \\ \min_{a \in action(s)} minimax(result(s, a)) & \text{if } player(s) = MIN \end{cases}$$

Properties of Minimax

- Complete? Yes, if tree is finite (chess has specific rules for this)
- Optimal? Yes, against an optimal opponent. Otherwise??
- Time? $O(b^m)$
- Space? $O(bm)$ (depth-first exploration)
For chess, $b \approx 35$, $m \approx 100$ for “reasonable” games
→ exact solution completely infeasible
But do we need to explore every path?


Minimax Properties



Optimal against a perfect player. Otherwise?

Resources Limits

- **Problem:** In realistic games, cannot search to leaves!
- **Solution:** Depth-limited search
 - Instead, search only to a limited depth in the tree
 - Replace terminal utilities with an evaluation function for non-terminal positions
- **Example:**
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program
- **Guarantee of optimal play is gone**
- **More plies makes a BIG difference**
- **Use iterative deepening for an anytime algorithm**




Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation



Evaluation Function

- Evaluation functions score non-terminals in depth-limited search




- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.


Game Tree Pruning



Minimax Example




Minimax Pruning



Alpha-Beta Pruning

- General configuration (MIN version)
 - We're computing the MIN-VALUE at some node n
 - We're looping over n 's children
 - n 's estimate of the childrens' min is dropping
 - Who cares about n 's value? MAX
 - Let a be the best value that MAX can get at any choice point along the current path from the root
 - If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)
- MAX version is symmetric



Alpha-Beta Pruning - Intuition

Two values:

- α = value of best choice so far for MAX (highest-value)
- β = value of best choice so far for MIN (lowest-value)
- Each node keeps track of its $[\alpha, \beta]$ values

α - β pruning algorithm

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in ACTIONS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) **returns** a *utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for each *a* **in** ACTIONS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** *v*

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return *v*

function MIN-VALUE(*state*, α , β) **returns** a *utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

for each *a* **in** ACTIONS(*state*) **do**


$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** *v*

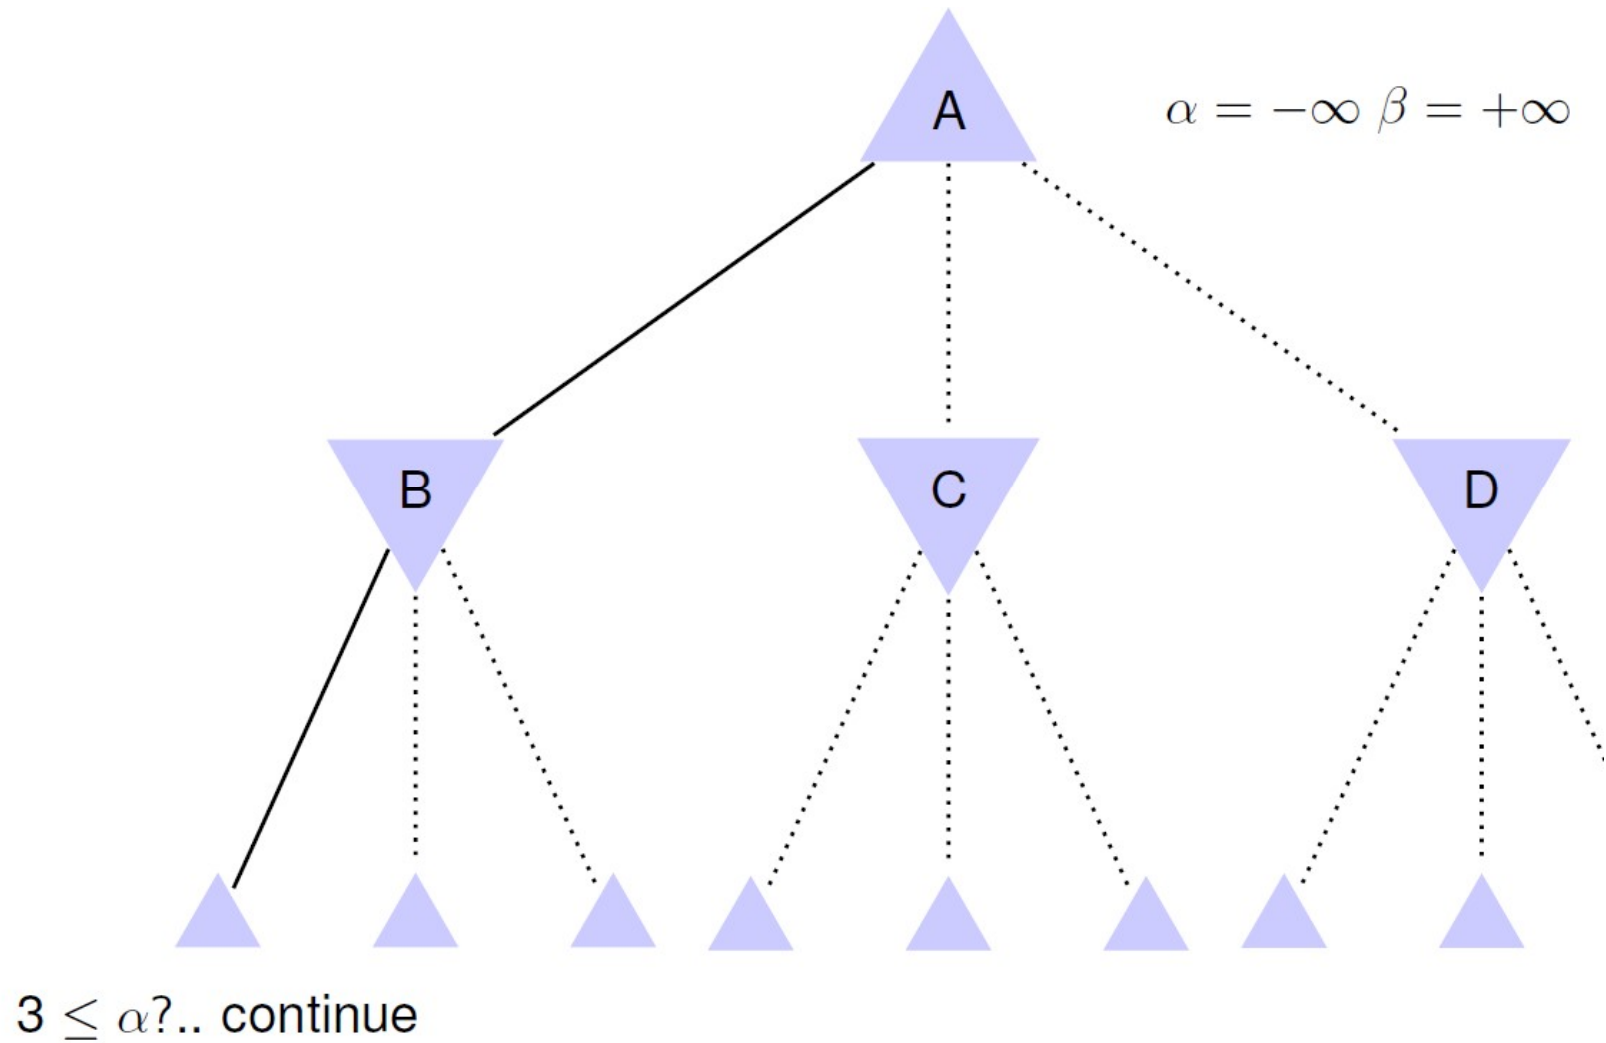
$\beta \leftarrow \text{MIN}(\beta, v)$

return *v*


α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example



α - β pruning example



α - β pruning example



Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - So the most naïve version won't let you do action selection
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of, e.g. chess, is still hopeless...